

# Anthropic Company Profile

Part 1 of 3 - PDF section, 20 pages. Prepared as a concise business and technical profile of Anthropic, the company behind Claude.

| Topic     | Details                                              |
|-----------|------------------------------------------------------|
| Company   | Anthropic, PBC                                       |
| Known for | Claude AI assistant and model family                 |
| Founded   | 2021                                                 |
| Focus     | Frontier AI, enterprise AI, safety, interpretability |

## PDF Page 1 - Executive Snapshot

Anthropic is one of the central companies in the frontier AI market. Its brand is built around Claude, a family of large language models used for writing, analysis, coding, enterprise automation, and software development assistance.

Anthropic, PBC is an American artificial intelligence company headquartered in San Francisco. It is best known for the Claude family of large language models and for positioning AI safety, interpretability, and steerability as central engineering goals.

The company differentiates itself by pairing commercial model deployment with an explicit safety research agenda. This framing makes Anthropic both a product company and a governance-oriented AI lab.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 2 - Origins and Founding Thesis

The company was founded in January 2021 by former OpenAI employees including Dario Amodei, Daniela Amodei, Jared Kaplan, Jack Clark, Chris Olah, Ben Mann, Sam McCandlish, and Tom Brown.

The founding thesis was that highly capable AI systems should be developed with safety, alignment, and interpretability research integrated from the beginning rather than added after deployment.

This early identity still shapes Anthropic's messaging: the company repeatedly emphasizes reliability, harmlessness, honesty, and controllability as part of its model development process.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 3 - Leadership and Organization

Dario Amodei is the CEO and Daniela Amodei is the President. The company has also been associated with prominent AI researchers focused on alignment and interpretability.

The leadership narrative matters because Anthropic emerged from a group of researchers who had already worked on scaling laws, transformer interpretability, and alignment methods before forming the company.

Its public-benefit-corporation structure signals that the company wants to balance shareholder value with wider public-interest commitments, although it still operates in a highly competitive commercial market.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 4 - Claude as the Core Product

The Claude product family includes chat, API access, Claude Code for software development workflows, and enterprise products such as Claude for Work and specialized workflow environments.

Claude is offered to consumers directly and to developers through APIs and cloud marketplaces. For enterprises, the value proposition is not only chat but integration: secure access, administrative controls, long-context reasoning, coding assistance, and domain workflow support.

Claude competes with other major assistants by focusing on long-form reasoning, professional writing, analysis, and developer productivity.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 5 - Business Model

Anthropic earns revenue through subscriptions, API usage, cloud-platform distribution, enterprise contracts, and coding-tool adoption. Usage-based API pricing is especially important because customers can embed Claude into their own applications and internal tools.

In the same May 2026 announcement, Anthropic said run-rate revenue had crossed \$47 billion earlier that month, reflecting rapid enterprise and developer adoption.

The business model is capital intensive because frontier models require massive compute for training, serving, evaluation, and safety testing.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 6 - Funding Scale

On May 28, 2026, Anthropic announced a \$65 billion Series H round at a \$965 billion post-money valuation, led by Altimeter Capital, Dragoneer, Greenoaks, and Sequoia Capital.

Anthropic said the funding would advance safety and interpretability research, expand compute capacity, and scale products and partnerships. That framing shows how tightly the company's product growth is tied to infrastructure spending.

Large private funding rounds also place pressure on Anthropic to sustain revenue growth, win enterprise customers, and maintain credibility in both safety and performance.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 7 - Compute Strategy

On Apr. 20, 2026, Anthropic and Amazon announced an expanded collaboration for up to 5 gigawatts of compute capacity for training and deploying Claude, with AWS remaining Anthropic's primary cloud provider and training partner.

Anthropic said more than 100,000 customers run Claude on Amazon Bedrock, and that the two companies launched Project Rainier, described as one of the largest compute clusters in the world.

Compute has become a strategic asset for Anthropic. Training and serving frontier models requires specialized accelerators, power availability, networking, storage, and teams capable of operating large distributed systems.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 8 - Multi-Cloud Distribution

Anthropic has emphasized that Claude is available across major cloud platforms including AWS, Google Cloud, and Microsoft Azure, supporting enterprise purchasing and governance requirements.

This multi-channel distribution strategy helps Anthropic reach customers that already standardize on a particular cloud vendor. It also reduces friction for security, procurement, compliance, and billing teams.

At the same time, cloud partnerships create strategic dependencies because model access, inference capacity, and customer acquisition are connected to infrastructure providers.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.



## PDF Page 9 - Constitutional AI

Anthropic's Constitutional AI approach uses a written set of principles to guide model self-critique, revision, and reinforcement learning from AI feedback, aiming to improve harmlessness without making the assistant overly evasive.

The method contains a supervised phase in which model outputs are critiqued and revised, and a reinforcement-learning phase using AI feedback. The key idea is making principles explicit so that alignment behavior is more inspectable.

Constitutional AI became central to Anthropic's identity because it offered a practical story for scaling safety feedback without relying only on human raters for every harmful scenario.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 10 - Responsible Scaling Policy

Anthropic's Responsible Scaling Policy defines AI Safety Levels for matching safeguards to model capability and catastrophic-risk potential, with higher levels requiring stronger safety and security measures.

The policy focuses on catastrophic risks such as misuse for dangerous capabilities or autonomous behavior that could cause large-scale harm. It is designed to evolve as evaluation science and model capabilities evolve.

For customers and policymakers, the RSP functions as a public accountability document: it explains what Anthropic says it will check before building or releasing more capable systems.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 11 - Interpretability Research

Anthropic is widely associated with mechanistic interpretability: research that tries to understand how neural networks represent concepts, features, and behaviors internally.

Interpretability is strategically important because it may help developers detect deception, bias, dangerous capability, or hidden failure modes before deployment.

This research area is still early, but it supports Anthropic's claim that frontier AI should be inspectable rather than treated as a black box.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 12 - Enterprise Adoption

Enterprise adoption is central to Anthropic's growth story. Claude is used for coding, document drafting, data analysis, customer support, internal workflow automation, legal-review support, and knowledge-base assistants.

The company's economic value is strongest when Claude becomes embedded into daily workflows rather than used as a standalone experimental chatbot.

Enterprise customers care about reliability, compliance, latency, budget controls, privacy promises, and vendor support.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 13 - Claude Code

Claude Code is Anthropic's developer-oriented product for software engineering workflows. It targets tasks such as reading codebases, editing files, debugging, writing tests, creating pull requests, and automating repetitive development work.

Developer adoption matters because programming is one of the clearest near-term areas where LLMs can produce measurable productivity gains.

Claude Code also strengthens API and subscription demand because coding agents can generate high-volume tool calls and long-running sessions.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 14 - Competition

Anthropic competes against OpenAI, Google DeepMind/Gemini, Meta, xAI, Mistral, Cohere, and other labs. Competition occurs across model quality, price, latency, context length, tool use, safety controls, and enterprise integration.

The market is not purely consumer-chat driven. For Anthropic, professional work, coding, developer platforms, and enterprise procurement are especially important battlegrounds.

Because models improve quickly, competitive position can change within months.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 15 - Governance and Policy Positioning

Anthropic has taken a comparatively visible role in AI safety policy discussions. Its public documents discuss evals, risk thresholds, red-teaming, and safety levels.

This positioning can increase trust with some enterprises and policymakers, but it can also attract scrutiny from people who worry about regulatory capture, concentration of power, or the feasibility of self-governance.

The important point is that governance is not separate from Anthropic's strategy; it is part of the company's brand and market differentiation.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 16 - Risks and Challenges

Anthropic faces several major risks: compute costs, competition, model misuse, regulatory friction, talent retention, copyright and training-data disputes, infrastructure dependency, and expectations created by very high valuations.

The company also has to balance helpfulness and safety. If Claude is too restrictive, customers may prefer competitors; if it is too permissive, Anthropic's safety-first reputation can be damaged.

Managing that balance is one of the company's defining execution challenges.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.



## PDF Page 17 - Product Design Philosophy

Claude's design philosophy emphasizes being helpful, harmless, and honest. In user-facing terms, this means producing useful answers while refusing or redirecting unsafe requests and acknowledging uncertainty when appropriate.

For enterprise users, the philosophy also appears in admin tooling, model documentation, and deployment controls.

The practical challenge is consistency: safety behavior must remain stable across languages, domains, tools, and long conversations.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 18 - Strategic Importance of Partners

Partners such as Amazon, Google, Microsoft, semiconductor suppliers, and enterprise software platforms are essential to Anthropic's scaling path.

The company needs cloud capacity to train and deploy models; it needs distribution partners to reach buyers; and it needs chip and memory supply chains to keep pace with demand.

These partnerships can be strengths, but they also create negotiation and dependency risks.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 19 - Outlook

Anthropic's outlook depends on whether Claude can continue to deliver high-quality reasoning and coding performance while maintaining trust, reliability, and safety credibility.

If enterprise AI becomes deeply embedded in everyday work, Anthropic could remain one of the dominant vendors. If models commoditize or compute costs overwhelm margins, the company may face pressure despite large revenue growth.

The next phase of competition will likely focus on autonomous agents, workflow products, model efficiency, and enterprise governance.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.

## PDF Page 20 - PDF Section Summary

Anthropic is best understood as a frontier AI company, an enterprise software vendor, and a safety-focused research lab at the same time.

Its rise has been driven by Claude adoption, large funding rounds, major cloud partnerships, and a strong narrative around safe and steerable AI.

The company's success will depend on execution across product quality, infrastructure, safety governance, and customer trust.

Source note: Anthropic, "Anthropic raises \$65B in Series H funding at \$965B post-money valuation," May 28, 2026. | Anthropic, "Anthropic and Amazon expand collaboration for up to 5 gigawatts of new compute," Apr. 20, 2026. | Anthropic, "Constitutional AI: Harmlessness from AI Feedback," Dec. 15, 2022. | Anthropic, "Responsible Scaling Policy," latest public policy page accessed July 2026.